

Phoneme-AI REST API Developer SOP

Production-ready SOP for developers to integrate Phoneme-AI into enterprise systems, automation, chatbots, and AI workflows.

1. Base URL

Primary Server URL: `http://10.100.60.121:11434`

Base URL:
`http://10.100.60.121:11434`

API Format:
`http://10.100.60.121:11434/api/<api-name>`

2. Available Production Models & Recommended Use Cases

Model	Best Use Case
phi3:latest	Fast testing, quick Q&A, light automation
gemma4:e4b	Balanced business tasks, backend APIs, docs
qwen3.6:27b	Deep reasoning, advanced coding, architecture design
qwen2.5:latest	Best all-rounder for coding, chatbots, daily enterprise

3. Important APIs with Description, Use Cases & Code Snippets

List Installed Models (/api/tags)

Endpoint: `http://10.100.60.121:11434/api/tags`

Description: Returns installed AI models.

Use Case: Check model availability before development.

```
import requests
r = requests.get("http://10.100.60.121:11434/api/tags")
print(r.json())
```

Generate Text (/api/generate)

Endpoint: `http://10.100.60.121:11434/api/generate`

Description: Single prompt direct response.

Use Case: Summaries, SQL, automation, reports.

```
import requests
payload = {
    "model": "gemma4:e4b",
    "prompt": "Explain JWT simply",
    "stream": False
}
r = requests.post("http://10.100.60.121:11434/api/generate", json=payload)
print(r.json()["response"])
```

Chat API (/api/chat)

Endpoint: `http://10.100.60.121:11434/api/chat`

Description: Multi-turn role-based chat.

Use Case: Chatbots, support, enterprise assistant.

```
import requests
payload = {
    "model": "qwen2.5:latest",
    "messages": [
        {"role": "system", "content": "You are Phoneme-AI enterprise assistant"},
        {"role": "user", "content": "How does Redis work?"}
    ],
    "stream": False
}
r = requests.post("http://10.100.60.121:11434/api/chat", json=payload)
print(r.json()["message"]["content"])
```

Pull Model (/api/pull)

Endpoint: `http://10.100.60.121:11434/api/pull`

Description: Downloads new AI model.

Use Case: Add new models for deployment.

```
import requests
requests.post("http://10.100.60.121:11434/api/pull",
              json={"name": "llama3:8b"})
```

Show Model Details (/api/show)

Endpoint: `http://10.100.60.121:11434/api/show`

Description: Shows model metadata.

Use Case: Audit architecture and suitability.

```
import requests
r = requests.post("http://10.100.60.121:11434/api/show",
                  json={"name": "gemma4:e4b"})
print(r.json())
```

Embeddings (/api/embeddings)

Endpoint: `http://10.100.60.121:11434/api/embeddings`

Description: Converts text to vectors.

Use Case: RAG, search, vector DB.

```
import requests
payload = {
    "model": "nomic-embed-text",
    "prompt": "What is JWT?"
}
r = requests.post("http://10.100.60.121:11434/api/embeddings", json=payload)
print(r.json())
```

4. API Parameter Reference

<code>model</code>	-> AI model
<code>prompt</code>	-> Single-shot input
<code>messages</code>	-> Chat history
<code>stream</code>	-> Live output
<code>temperature</code>	-> Creativity
<code>top_p</code>	-> Variability
<code>timeout</code>	-> Request timeout

5. Troubleshooting

Connection Refused	-> Verify server
404 Model Error	-> Check /api/tags
Slow Response	-> Use phi3
Timeout	-> Increase timeout

Final Recommendation

Default: qwen2.5:latest for most enterprise tasks. Use phi3 for speed, gemma4 for balanced workflows, qwen3.6 for advanced reasoning.